

A Stylometric Application of Large Language Models

Anonymous ACL submission

Abstract

In this paper we show that large language models (LLMs) can be used to distinguish the writings of different authors. Specifically, an individual model, trained on the works of one author, will predict held-out text from that author more accurately than held-out text from other authors. We suggest that, in this way, a model trained on one author’s works embodies the unique writing style of that author. We first demonstrate our approach on books written by eight different (known) authors. We also use this approach to confirm R. P. Thompson’s authorship of the well-studied 15th book of the *Oz* series, originally attributed to F. L. Baum.

1 Introduction

In this paper we introduce *predictive comparison*, a new LLM-based relative stylometric measure. It derives from a simple idea, that if an LLM can be fine-tuned to write like—i.e., in the style of—a given author by training on the work of an author (e.g., Mikros, 2025), then the degree to which such a fine-tuned model can predict another author’s work could be a measure of stylistic similarity. In this paper we show, using a small set of authors and their works, that this thesis is borne out. This in turn suggests a notion of stylometric distance derived from the cross-entropy loss assigned to held-out texts by models trained on known works of different authors. Lastly, this further suggest a literary authentication tool that would assign an unknown or contested work to the model which predictive comparison generates the smallest loss. We use this on a well-known and once contested 15th book in the *Oz* series, confirming what is now accepted attribution. We believe this approach could be of use in considering questions of autho-

rial influence (Mosteller and Wallace, 1963, 1984; Binongo, 2003; Juola, 2008) and stylistic evolution.

2 Methods

In this section, we outline our methodology for identifying stylometric signatures using large language models. For each selected author, we train a GPT-2 model (Radford et al., 2019) on that author’s corpus. We then use the trained model to compute the prediction loss on some held-out texts from both the target author and each of the other authors in the dataset. By comparing these losses, we assess whether the model captures author-specific stylistic patterns: a model trained on a given author should exhibit lower loss when predicting that author’s own texts compared to those of others.

2.1 Data and Preprocessing

We consider a dataset comprising books by eight authors: Jane Austen, L. Frank Baum, Charles Dickens, F. Scott Fitzgerald, Herman Melville, Rosemary Plumbly Thompson, Mark Twain, and H. G. Wells. We selected these authors because their writings are well-represented in Project Gutenberg, are all in the public domain, and are written in English—eliminating any potential confounds due to translation. For each book, we pre-process the text by stripping Project Gutenberg metadata, publisher information, illustration tags, transcriber notes, prefaces, tables of contents, and chapter headings. We standardize whitespace, remove non-ASCII characters, and lowercase all alphabetic characters. Basic statistics on token lengths and the full list of books used are provided in the Appendix.

To construct a training data for each author, we randomly select one book to hold out for evalua-

tion and train their model using the remaining books. To ensure fair comparisons across authors, we standardize the number of training tokens per author. Specifically, we truncate each author’s corpus so that every model is trained on an equal number of tokens. This token budget is determined by removing the longest book from each author’s set and then taking the smallest remaining total token count across each author’s remaining books. For our dataset, this yields a training token budget of 643,041 tokens.

To construct a truncated corpus of 643,041 tokens for each author, we sample one contiguous subsequence from each book in their training corpus (i.e., remaining books after holding out a to-be-evaluated book). The length of the sub-sequence sampled from book i is proportional to its original length, computed as:

$$\text{length}_i = 643,041 \times \frac{\text{tokens in book } i}{\text{total tokens in corpus}}.$$

The starting position of each sub-sequence is chosen uniformly at random, ensuring the sample fits within the book’s bounds. Finally, we shuffle and then concatenate the sampled sub-sequences from each book, resulting in a single 643,041-token training corpus for each author. This process is repeated for each of 10 random seeds, yielding 10 different training corpora for each author.

2.2 Model Architecture, Training, and Evaluation

For each author, we train GPT-2 language models (Radford et al., 2019) from scratch using the GPT2LMHeadModel class from the Hugging Face Transformers library with custom architecture settings: a context window of 1024 tokens, an embedding dimension of 128, 8 transformer layers, and 8 attention heads per layer. We fit each model using the AdamW optimizer (Loshchilov and Hutter, 2019) with a learning rate of 5×10^{-5} to minimize the cross-entropy loss on the training data. We train models using a causal language modeling objective, whereby the model predicts the next token in the sequence given all of the previous tokens in the same batch.

We construct training samples by shuffling, concatenating, and then sampling 1024-token chunks

from the truncated corpus for the given author (constructed as described above, using contiguous subsequences selected from all but one of their books). Each training epoch consists of 40 batches, each containing 16 sequences of 1024 tokens. This results in a total of 655,360 tokens per epoch. We continue training until the cross-entropy loss fell to 3.0 or lower. (We decided on this threshold after taking random draws from the models trained on Baum’s and Thompson’s *Oz* books and manually inspecting the quality of the resulting samples.) Training to a fixed loss threshold (e.g., as opposed to training for a fixed number of epochs) enables us to fairly compare model performance across authors, which is a critical component of our stylometric analyses.

We evaluate the models using the held-out book from the corresponding author. We sample 1024-token chunks from the held-out book, using a sliding window approach to ensure that each token in the evaluation set contributes equally to the computed loss. We repeat the full process (of selecting a held-out book at random and training the model using randomly selected samples from the remaining books) using 10 different random seeds. This approach enables us to assess the robustness of our results and to ensure that the models are not overfitting to a specific book or random sample.

3 Results

3.1 Predictive Comparison Testing of Eight Classic Authors

We carried out predictive comparison testing on eight classic authors (see Sec. 2.1). The top-left sub-panel of Figure 1A (labeled “Train”) shows the average training loss for each author’s model, computed over 10 random seeds. Training losses are comparable across models, indicating that the models are trained to similar levels of performance. The other sub-panels of Figure 1A show the average predictive (cross-entropy) loss, for each author’s model, on held-out texts from each author. For every author’s held-out text, the model trained on the same author produces the lowest loss, indicating a clear preference for its own author’s stylistic patterns. As shown in Figure 1B, across every author we considered, and for every random seed, models trained and tested on the same author always yield smaller losses than

models trained on one author and tested on another (i.e., for each author, the highest black dot is always lower than the lowest gray dot in the Panel). Indeed, we achieve perfect (100%) classification accuracy for matching authors with held-out texts by labeling the held-out text according to which model produces the smallest loss.

We also wondered how many training epochs were required for the models to reliably distinguish author styles. To investigate this, we compared the distributions (across random seeds) of average cross-entropy losses for each author’s model computed for held-out text from the *same* author versus for held-out text from *other* authors. Figure 1D displays the t -values from paired t -tests comparing these same vs. other loss distributions for each of the first 500 training epochs. For all authors except Twain, the t -tests yielded p -values below 0.001 after just one epoch, indicating that the models rapidly acquire author-specific stylometric patterns. For Twain, this threshold is crossed at epoch 47. Figure 1E shows the average t -values across all eight authors as a function of the number of training epochs (final epoch: $t(9) = 20.723, p = 6.6 \times 10^{-9}$). This latter plot provides an estimate of the performance we might expect to see in the general case (e.g., across a larger set of authors). Table 1 summarizes the results of the t -tests for each author’s model after 500 training epochs.

Model	t -stat	df	p -value
Baum	16.96	10.49	5.78×10^{-9}
Thompson	21.50	13.60	6.84×10^{-12}
Austen	47.29	54.75	4.38×10^{-46}
Dickens	18.36	27.36	6.52×10^{-17}
Fitzgerald	26.03	22.66	2.22×10^{-18}
Melville	24.15	45.15	1.87×10^{-27}
Twain	20.13	12.22	9.67×10^{-11}
Wells	35.17	26.33	1.16×10^{-23}

Table 1: Each row displays the results of a paired t -test comparing the average loss values assigned by each author’s model (after 500 training epochs) to the author’s held-out text and to the other authors’ randomly sampled texts.

Despite achieving perfect classification accuracy, not all authors are equally distinctive. For example, we reasoned that authors with similar writing styles might be more confusable (i.e., yielding relatively small losses for models trained across different au-

thors). We computed the average loss for each author using the models trained on the other authors’ texts (Fig. 1F). Authors with similar writing styles (e.g., Baum and Thompson) yield relatively small losses when evaluated using models trained on the other author’s texts. In contrast, authors with more distinct writing styles (e.g., Austen and Thompson) yeidl relatively large lossess when evaluated using each other’s models. To illustrate these patterns, we also project the losses into a 3D space using multidimensional scaling (MDS) (Kruskal, 1964) applied to the pairwise correlations between rows of the loss matrix, excluding the diagonal entries (i.e., the losses obtained using each author’s model when applied to their own held-out text). We observed that Baum and Thompson (both authors of children’s literature); Twain and Melville (authors of adventure and exploration stories); Dickens and Austen (authors of romantic and domestic stories); and Fitzgerald and Wells (authors of speculative and modern fiction) each appear nearby in the MDS projection, suggesting that the models capture some stylistic features that are shared across similar authors.

3.1.1 Predicted Authorship of the 15th Oz Book

As a special case, we also consider the contested authorship of the 15th Oz book, which is widely believed to have been written by Ruth Plumly Thompson, but was originally attributed to L. Frank Baum (Binongo, 2003). We applied our predicting comparison testing to the 15th Oz book, using the models trained on Baum and Thompson’s undisputed Oz books. As shown in the bottom left sub-panel of Figure 1C, our approach yields lower loss for the Thompson-trained model than for the Baum-trained model, indicating that the contested book is indeed more similar to Thompson’s writing style than to Baum’s. We also applied both models to a non-Oz book by Baum (bottom center) and Thompson (bottom right). Our approach yields lower losses for the correct author in each case, demonstrating that PCT is robust to thematic differences from the same author.

3.2 Stylometric Distance

Predictive comparison suggests a natural notion of distance between authorial styles. Let $L_i(j)$ denote the average loss of a work of author j for

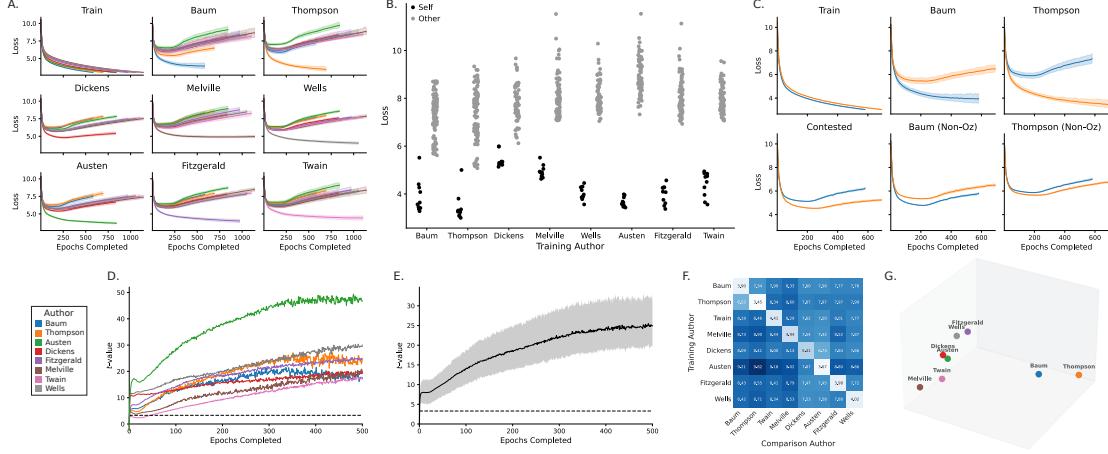


Figure 1: **A.** Average cross-entropy loss on *Training* data and held-out test data from each author, plotted as a function of the number of training epochs. Each color denotes a model trained on a single author’s work. Error ribbons denote bootstrap-estimated 95% confidence intervals over 10 random seeds. **B.** Cross-entropy loss assigned to held-out test data by each author’s model (*x*-axis). Held-out test data is either from the *same* author (block) or from *other* authors (gray). Each dot denotes the average loss (across all samples) for a single random seed. **C.** The top sub-panels replicate the Baum (blue) and Thompson (orange) results from Panel A. The bottom sub-panels show the cross-entropy loss assigned to a held-out text whose authorship is contested (lower left), to a held-out non-*Oz* text by Baum (lower center), and to a held-out non-*Oz* text by Thompson (lower right). Error ribbons denote bootstrap-estimated 95% confidence intervals over 10 random seeds. **D.** Each curve denotes, as a function of the number of training epochs, the the *t*-statistic from a *t*-test comparing the distribution of losses (across random seeds) assigned to held-out texts from the given author (color) versus held-out texts from all other authors. **E.** The average *t*-statistic across all eight authors, as a function of the number of training epochs. Error ribbons denote bootstrap-estimated 95% confidence intervals across authors. **F.** The average cross-entropy loss assigned by models trained on each author’s writing (row) to held-out texts from each author (column). **G.** Three-dimensional MDS projection of the average cross-entropy loss matrix shown in Panel F.

a model trained on author i (the i, j -entry of the heatmap/average loss matrix in Figure 1F). Let $\overline{L}_i(j) = L_i(j) - L_i(i)$, normalizing the entries by subtracting the native author baseline. Then define the LLM-based *stylometric distance*, $d(i, j) = \frac{1}{2} \left(\overline{L}_i(j) + \overline{L}_j(i) \right)$. Figure 1G is a visualization of the relative “distances” among our author set.

Conclusions

Just as prior work has shown that it is possible to fine-tune LLMs to *write* in the “style” or “voice” of a given author (see e.g., Mikros, 2025), our work shows that LLMs may also be used to predict authorship and measuring the stylistic distances between different authors. We note that our approach is broadly similar to that of (Rezaei, 2025), who examined sentence level information using LLMs. Our work differs in scale (we use entire books rather than individual sentences) and in our reliance solely on cross-entropy loss as a measure of stylistic dis-

tance. We suggest that our approach holds promise as a new technique for machine reading approaches to text-based disciplines (Moretti, 2017, 2000; Holmes, 1998) and the practices of cultural analytics (Underwood et al., 2013).

Limitations

The main limitations of this paper are (1) the lack of breadth of experiments as well as (2) the oft-acknowledged opacity of the LLM. The results in this paper serve as a proof-of-concept for the idea of using the structure of a bespoke trained LLM as a stylometric engine. Only a handful of examples have been tested, but the results on a classic stylometric test are promising. Further testing is needed to understand what kinds of writing features are being picked up by the LLM. Finally, deploying this idea at scale would require training one model for every writer of interest, a task that would require significant computational and textual resources.

241
242
243
244
245
246
247

248

249
250
251
252
253
254
255
256
257
258
259

260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278

References

- 280 José Nilo G. Binongo. 2003. Who wrote the 15th book of
 281 Oz? An application of multivariate analysis to author-
 282 ship attribution. *CHANCE*, 16(2):9–17.
- 283 David I Holmes. 1998. The evolution of stylometry in
 284 humanities scholarship. *Literary and Linguistic Com-*
 285 *puting*, 13(3):111—117.
- 286 Patrick Juola. 2008. Authorship attribution. *Foundations*
 287 *and Trends in Information Retrieval*, 1(3):233–334.
- 288 Joseph B Kruskal. 1964. Multidimensional scaling by
 289 optimizing goodness of fit to a nonmetric hypothesis.
 290 *Psychometrika*, 29(1):1–27.
- 291 Ilya Loshchilov and Frank Hutter. 2019. De-
 292 coupled weight decay regularization. *Preprint*,
 293 arXiv:1711.05101.
- 294 George Mikros. 2025. Beyond the surface: stylometric
 295 analysis of GPT-4o’s capacity for literary style imi-
 296 tation. *Digital Scholarship in the Humanities*, page
 297 fqaf035.
- 298 Franco Moretti. 2000. Conjectures on world literature.
 299 *New Left Review*, 1:54–68.
- 300 Franco Moretti. 2017. *Graphs, Maps, Trees: Abstract*
 301 *Models for Literary History*. Verso Books, Brookly,
 302 NY.
- 303 Frederick Mosteller and David L. Wallace. 1963. Infer-
 304 ence in an authorship problem. *Journal of the Ameri-*
 305 *can Statistical Association*, 58(302):275–309.
- 306 Frederick Mosteller and David L. Wallace. 1984. *Applied*
 307 *Bayesian and Classical Inference: The Case of the*
 308 *Federalist Papers*. Addison-Wesley, Reading, MA.
- 309 Alec Radford, Jeffrey Wu, Rewon Child, David Luan,
 310 Dario Amodei, Ilya Sutskever, and 1 others. 2019.
 311 Language models are unsupervised multitask learners.
 312 *OpenAI blog*, 1(8):9.
- 313 Mosab Rezaei. 2025. Detecting, generating, and eval-
 314 uating in the writing style of different authors. In
 315 *Proceedings of the 2025 Conference of the Nations*
 316 *of the Americas Chapter of the Association for Com-*
 317 *putational Linguistics: Human Language Technolo-*
 318 *gies (Volume 4: Student Research Workshop)*, pages
 319 485–491, Albuquerque, USA. Association for Compu-
 320 *tational Linguistics*.
- 321 Ted Underwood, Michael L. Black, Loretta Auvil, and
 322 Boris Capitanu. 2013. Mapping mutable genres in
 323 structurally complex volumes. In *2013 IEEE Interna-*
 324 *tional Conference on Big Data*, page 95–103. IEEE.

Appendix: Authors, Books, Tokens

Charles Dickens	Tokens	Herman Melville	Tokens
A Christmas Carol	38,906	I and My Chimney	15,341
Oliver Twist	216,100	Bartleby, the Scrivener	19,112
The Old Curiosity Shop	285,895	Israel Potter	88,570
Bleak House	471,630	Omoo	134,628
Dombey and Son	482,161	Mardi, Vol. II	150,347
David Copperfield	479,387	The Confidence-Man	129,059
A Tale of Two Cities	181,593	White Jacket	190,577
Nicholas Nickleby	446,457	Mardi, Vol. I	132,358
American Notes	129,214	Moby-Dick	285,066
The Pickwick Papers	432,546	Typee	114,239
Great Expectations	244,897		
Martin Chuzzlewit	455,995		
Little Dorrit	449,230		
Hard Times	142,759		
Total	4,456,770	Total	1,259,297

L. Frank Baum	Tokens	Ruth Plumly Thompson	Tokens
Ozma of Oz	52,039	The Giant Horse of Oz	51,036
Dorothy and the Wizard in Oz	53,849	The Cowardly Lion of Oz	61,666
Tik-Tok of Oz	63,781	Handy Mandy in Oz	44,778
The Road to Oz	52,866	The Gnome King of Oz	51,687
The Magic of Oz	51,166	Grampa in Oz	55,169
The Patchwork Girl of Oz	75,703	Captain Salt in Oz	61,797
The Wonderful Wizard of Oz	49,686	Ozoplaning with the Wizard of Oz	50,660
The Lost Princess of Oz	60,418	The Wishing Horse of Oz	59,490
The Emerald City of Oz	70,781	The Lost King of Oz	58,105
The Tin Woodman of Oz	57,338	The Hungry Tiger of Oz	53,543
Rinkitink in Oz	62,241	The Silver Princess in Oz	47,964
The Marvelous Land of Oz	54,733	Kabumpo in Oz	62,693
Glinda of Oz	51,218	Jack Pumpkinhead of Oz	49,661
The Scarecrow of Oz	59,593		
Total	815,412	Total	708,249

Austen	Tokens	Twain	Tokens
Sense And Sensibility	153,718	Adventures Of Huckleberry Finn	147,655
Mansfield Park	201,611	A Connecticut Yankee In King Arthur'S Court	150,327
Lady Susan	29,043	Roughing It	208,545
Northanger Abbey	98,090	The Innocents Abroad	246,321
Emma	207,830	The Adventures Of Tom Sawyer, Complete	95,059
Pride And Prejudice	157,777	The Prince And The Pauper	88,409
Persuasion	106,027		
Total	954,096	Total	936,316

Fitzgerald	Tokens	Wells	Tokens
The Beautiful And Damned	168,147	The Red Room	4,944
Flappers And Philosophers	84,707	The First Men In The Moon	87,615
This Side Of Paradise	100,796	The Island Of Doctor Moreau	55,967
All The Sad Young Men	85,411	The Open Conspiracy	40,271
Tales Of The Jazz Age	109,997	A Modern Utopia	105,810
The Pat Hobby Stories	51,069	The Sleeper Awakes	98,228
The Great Gatsby	65,136	The New Machiavelli	185,158
Tender Is The Night	145,925	The War Of The Worlds	75,727
		Tales Of Space And Time	94,711
		The Invisible Man: A Grotesque Romance	65,584
		The Time Machine	40,184
		The World Set Free	80,518
Total	811,188	Total	934,717